

```
from google.colab import drive
drive.mount('/content/drive')
```

```
import numpy as np
import pandas as pd
```

```
path="/content/drive/MyDrive/Churn_Modelling.csv"
df=pd.read_csv(path)
df.describe
```

```
<bound method NDFrame.describe of
CreditScore Geography Gender Age \
0          1    15634602   Hargrave      619    France  Female    42
1          2    15647311     Hill      608     Spain  Female    41
2          3    15619304     Onio      502    France  Female    42
3          4    15701354     Boni      699    France  Female    39
4          5    15737888  Mitchell      850     Spain  Female    43
...      ...      ...      ...      ...      ...      ...
9995      9996    15606229   Obijiaku      771    France    Male    39
9996      9997    15569892  Johnstone      516    France    Male    35
9997      9998    15584532      Liu      709    France  Female    36
9998      9999    15682355  Sabbatini      772   Germany    Male    42
9999     10000    15628319     Walker      792    France  Female    28
```

```
      Tenure      Balance  NumOfProducts  HasCrCard  IsActiveMember \
0          2         0.00             1           1             1
1          1    83807.86             1           0             1
2          8   159660.80             3           1             0
3          1         0.00             2           0             0
4          2   125510.82             1           1             1
...      ...      ...      ...      ...      ...
9995        5         0.00             2           1             0
9996       10    57369.61             1           1             1
9997        7         0.00             1           0             1
9998        3    75075.31             2           1             0
9999        4   130142.79             1           1             0
```

```
      EstimatedSalary  Exited
0          101348.88       1
1          112542.58       0
2          113931.57       1
3           93826.63       0
4           79084.10       0
...      ...      ...
9995          96270.64       0
9996       101699.77       0
9997          42085.58       1
9998          92888.52       1
9999          38190.78       0
```

```
[10000 rows x 14 columns]>
```

```
import matplotlib.pyplot as plt

import seaborn as sns

%matplotlib inline


df[['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
    'Gender', 'Age', 'Tenure']].describe()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000




```
sns.histplot(df.Geography,kde=True)
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f57aa309350>



```
sns.countplot(df.Gender)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: P
FutureWarning
<matplotlib.axes._subplots.AxesSubplot at 0x7f5797ef58d0>
```




```
df[['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
    'Gender', 'Age', 'Tenure']].corr()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure
RowNumber	1.000000	0.004202	0.005840	0.000783	-0.006495
CustomerId	0.004202	1.000000	0.005308	0.009497	-0.014883
CreditScore	0.005840	0.005308	1.000000	-0.003965	0.000842
Age	0.000783	0.009497	-0.003965	1.000000	-0.009997
Tenure	-0.006495	-0.014883	0.000842	-0.009997	1.000000



```
sns.scatterplot(df.CreditScore, df.Age)
plt.ylim(0, 150)
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pas
FutureWarning
(0.0, 150.0)
```




Double-click (or enter) to edit

3 Multi - Variate Analysis

Double-click (or enter) to edit

```
sns.pairplot(data=df[['CustomerId', 'Surname', 'CreditScore', 'Geography', 'Gender', 'Age',
```

<seaborn.axisgrid.PairGrid at 0x7f57979ec4d0>



Double-click (or enter) to edit

4 Descriptive statistics

```
df.describe()
```

	RowNumber	CustomerId	CreditScore	Age	Tenure	Bala
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000	10000.000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800	76485.889
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174	62397.405
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000	0.000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000	0.000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000	97198.540
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000	127644.240
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000	250898.090



```
df.dtypes
```

```
RowNumber      int64
CustomerId      int64
Surname         object
CreditScore     int64
Geography       object
Gender          object
Age             int64
Tenure          int64
Balance         float64
NumOfProducts  int64
HasCrCard       int64
IsActiveMember  int64
EstimatedSalary float64
Exited          int64
dtype: object
```

```
df['Age'].mode()
```

```
0    37
dtype: int64
```

```
df["Age"].mean()
```

```
38.9218
```

```
round(df["Age"].mean(), 3)
```

```
38.922
```

```
df["Age"].median()
```

```
37.0
```

Double-click (or enter) to edit

5 Handling Missing Values

```
df.isna().any()
```

```
RowNumber      False
CustomerId      False
Surname         False
CreditScore     False
Geography       False
Gender          False
Age             False
Tenure          False
Balance         False
NumOfProducts  False
HasCrCard       False
IsActiveMember  False
EstimatedSalary False
Exited          False
dtype: bool
```

```
df.isnull().sum()
```

```

RowNumber      0
CustomerId     0
Surname         0
CreditScore    0
Geography      0
Gender         0
Age            0
Tenure         0
Balance        0
NumOfProducts 0
HasCrCard      0
IsActiveMember 0
EstimatedSalary 0
Exited         0
dtype: int64

```

df.isnull()

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenur
0	False	False	False	False	False	False	False	Fals
1	False	False	False	False	False	False	False	Fals
2	False	False	False	False	False	False	False	Fals
3	False	False	False	False	False	False	False	Fals
4	False	False	False	False	False	False	False	Fals
...
9995	False	False	False	False	False	False	False	Fals
9996	False	False	False	False	False	False	False	Fals
9997	False	False	False	False	False	False	False	Fals
9998	False	False	False	False	False	False	False	Fals
9999	False	False	False	False	False	False	False	Fals

10000 rows × 14 columns



df.notnull()

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure
0	True	True	True	True	True	True	True	True
1	True	True	True	True	True	True	True	True
2	True	True	True	True	True	True	True	True
3	True	True	True	True	True	True	True	True
4	True	True	True	True	True	True	True	True
...
9995	True	True	True	True	True	True	True	True
9996	True	True	True	True	True	True	True	True
9997	True	True	True	True	True	True	True	True
9998	True	True	True	True	True	True	True	True
9999	True	True	True	True	True	True	True	True

10000 rows × 14 columns




Double-click (or enter) to edit

6 Find the outliers and replace the outliers


```
sns.boxplot(x=df['Age'])
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f57978189d0>



```
sns.boxplot(x=df['Tenure'])
```


<matplotlib.axes._subplots.AxesSubplot at 0x7f5794c20a90>



Double-click (or enter) to edit

Double-click (or enter) to edit

7 Check for Categorical columns and perform encoding.

```
a=df.columns
```

```
b=df._get_numeric_data().columns
```

b

```
Index(['RowNumber', 'CustomerId', 'CreditScore', 'Age', 'Tenure', 'Balance',
      'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary',
      'Exited'],
      dtype='object')
```

```
list(set(a) - set(b))
```

```
['Gender', 'Geography', 'Surname']
```

Double-click (or enter) to edit


8 Split the data into dependent and independent variables.

Double-click (or enter) to edit

```
# x -Independent
# y -Dependent
x =df.drop('Exited',axis=1)
y=df['Exited']
```

x.head()

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age	Tenure	Ba
0	1	15634602	Hargrave	619	France	Female	42	2	
1	2	15647311	Hill	608	Spain	Female	41	1	838
2	3	15619304	Onio	502	France	Female	42	8	1590
3	4	15701354	Boni	699	France	Female	39	1	
4	5	15737888	Mitchell	850	Spain	Female	43	2	1251



y.head()

```
0    1
1    0
2    1
3    0
4    0
Name: Exited, dtype: int64
```

Double-click (or enter) to edit

9 Scale the independent variables

```
from sklearn import linear_model
```

```
from sklearn.preprocessing import StandardScaler
```

```
scale = StandardScaler()
```

```
x=df[['Age', 'Tenure']]
```

```
scaledx = scale.fit_transform(x)
```

```
print(scaledx)
```

```
[[ 0.29351742 -1.04175968]
 [ 0.19816383 -1.38753759]
 [ 0.29351742  1.03290776]
 ...
 [-0.27860412  0.68712986]
 [ 0.29351742 -0.69598177]
 [-1.04143285 -0.35020386]]
```

Double-click (or enter) to edit

10 Split the data into training and testing

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=0)
```

```
print('X Train shape:{},Y.Train SHape:{}'.format(x_train.shape,y_train.shape))
```

```
X Train shape:(8000, 2),Y.Train SHape:(8000,)
```

```
print('X Test Shape :{},Y Test SHape:{}'.format(x_test.shape,y_test.shape))
```

```
X Test Shape :(2000, 2),Y Test SHape:(2000,)
```

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 10:34 AM

